

ASSIGNMENT 2 OF 2 · REQUIRED COVER SHEET · DUE SUNDAY 13 SEPTEMBER 2026, 23:59 SGT

# A2 Team Self-Appraisal

Team ID: \_\_\_\_C-5\_\_\_\_ Section: \_\_C\_\_ Problem chosen (A / B): \_\_A\_\_

Members: \_\_Chen Baiyi, Qian Ruoqi, Fang Hioieng, Sood Vedansh, Yang Yichen &amp; Zhang Qizhi \_\_

Repository URL: \_\_https://github.com/vedanshsood/PE6201-A2-Insurance-Claims-Agent-Group-5\_\_ Date: 20/09/2026

ONE SHEET PER TEAM, not one per member. Complete it together, in one sitting, after the work is done — it takes about twenty minutes. CHECK IT INTO YOUR A2 SUBMISSION FOLDER alongside the rest of the submission, by the A2 deadline of Sunday 13 September 2026, 23:59 SGT; do not email it. This sheet is required and is not graded, but a missing self-appraisal is treated as an incomplete submission.

## 1 · Rate your team against Rubric 1

Circle or mark one band per criterion. Be honest rather than generous — a team that marks itself Excellent everywhere tells the marker nothing, and a large gap between your rating and the marker's is the single most useful thing on this page.

| Criterion                 | Weight | Excellent (4) | Proficient (3) | Developing (2) | Limited (1) |
|---------------------------|--------|---------------|----------------|----------------|-------------|
| Conceptual Understanding  | 25%    |               | ✓              |                |             |
| Technical Execution       | 30%    | ✓             |                |                |             |
| Reasoning & Justification | 25%    |               | ✓              |                |             |
| Communication & Clarity   | 20%    |               | ✓              |                |             |

## 2 · The one decision you would defend hardest

Name a single design decision — a tool you cut, the autonomy setting you chose, the dependency rule behind your parallel calls, the model you settled on — and say what evidence made it. Four or five sentences.

The design decision **we would defend hardest is keeping AUTONOMY="confirm" for the irreversible decision-letter action**. Our live evaluation showed that even our strongest tested configuration, GPT-OSS-20B, achieved **54/60 trials (90%)**, so incorrect or failed decisions were still possible. We also observed meaningful performance differences across model families, showing that reliable behaviour could not be assumed across models. As a result, we allow the agent to autonomously gather evidence and determine routing, but keep the final member-facing action behind human confirmation. This gives us the efficiency benefits of agentic automation without allowing a model error to directly become an irreversible claim decision.

## 3 · The one you are least sure about

Every team has one. Naming it costs you nothing here and it is often the most interesting thing in the submission.

We are least sure about **how well our V2 improvements generalise across different models**. The same system produced different results across model families, suggesting that improvements effective for one model may not transfer equally to others. This would require broader testing before deployment.

## 4 · Your headline numbers

Copy these from your repository. If any cell is empty, that is worth noticing before you submit rather than after.

|                                | Value              |                                 | Value                |
|--------------------------------|--------------------|---------------------------------|----------------------|
| Evaluation cases               | 42                 | Guardrail cases                 | 10                   |
| Negative cases                 | 9                  | Live runs per model (D4)        | 60                   |
| Pass rate (best model)         | 90.00%             | Pass rate on negatives only     | 77.80%               |
| Models in the battery          | 5                  | Median turns per run            | 4                    |
| The live model YOU ran (D5b)   | openai/gpt-oss-20b | Your pass rate on it            | 90.0% (54/60)        |
| Turns saved by parallel calls  | 84 (323 → 239)     | Cost per successful task        | \$0.760589           |
| Monthly cost at problem volume | \$6,084.71 / month | Break-even success rate         | 88.27%               |
| Tokens per call — v1 → v2      | 24.46 → 31.97      | Pass rate — v1 → v2, same model | <b>80.0% → 90.0%</b> |

Every member runs one live model, or the v1 pass — three models is the floor, N-1 for a team of N is the shape expected. Name YOURS. If you ran the v1 pass, write "v1 pass" and the model it sat on.

The last row is D2(b): the descriptor rewrite you measured. Two numbers each, before and after, on ONE model. An empty row here means the rewrite was never measured.

## 5 · Contribution

Who owned what. This should match CONTRIBUTIONS.md and the commit history. It is not used to scale marks — the team receives the team score — but it is one of the two signals the non-contribution clause relies on.

| Member       | Owned                                                     | Also contributed to                                               |
|--------------|-----------------------------------------------------------|-------------------------------------------------------------------|
| Sood Vedansh | Agent implementation, V2 integration & evaluation harness | Parallel execution, guardrails, live evaluation, failure analysis |
| Yang Yichen  | System reliability & design review                        | Dependency design, guardrail review, deployment limitations       |
| Zhang Qizhi  | Agent architecture & conceptual design                    | Agent justification, workflow and tool design                     |
| Chen Baiyi   | V1 baseline & tool-interface design                       | Tool descriptors, V1→V2 comparison, evaluation                    |
| Qian Ruoqi   | Evaluation & guardrail testing                            | Negative cases, cross-model result analysis                       |
| Fang Hioieng | Cost & deployment analysis                                | Cost model, sensitivity analysis, break-even analysis             |

Peer rating is separate from this sheet and comes AFTER it: NTU's Group Project peer-review function opens 13 September and closes Wednesday 16 September, 23:59 SGT. Each member completes it individually in NTULearn. It is a participation requirement.

## 6 · Declaration

Every member ticks every box. One sheet, signed by the team.

[ ✓ ] Every member of this team can explain every block of code we submit — what it does, and why it is there. (Block, not line: this is about function, not syntax.)

[ ✓ ] The pass rates, token counts, turn counts and costs we report are measurements we ran ourselves. None of them were supplied by an AI assistant.

[ ✓ ] Our submitted harness runs end to end on the scripted backend, with no key and no network.

[ ✓ ] We kept every fixture record we were given, unedited, and every case in our evaluation set carries a label. `check_my_data.py` passes on our data.

[ ✓ ] No part of this submission is drawn from any member's End-of-Course Project, and no part of it will be submitted as part of one.

[ ✓ ] This is our own work, and all sources, tools and assistance are attributed.



Signed on behalf of the team: \_\_\_\_\_ Date: \_20/09/2026\_

## How this is used

This sheet is not marked and cannot lower your grade. It is read alongside your submission for two reasons. First, how a team reads its own work is usually the fastest route to feedback that is actually useful. Second, where a team's self-rating and the marking diverge sharply, that gap gets an explicit comment — which is where the most useful feedback tends to sit.